

# **Semi-Supervised Learning for Fine-Grained Image Classification: Using Vision Transformers and Transfer Learning**

## **Authors**

Varun Sathyanarayanan, Harshith Ravi Kopparam

## **Abstract**

In this work, we present a semi-supervised learning approach for fine-grained image classification using a custom dataset containing 135 categories (15 plant and 120 dog categories). Our methodology leverages both labeled (9,854 images) and unlabeled (22,995 images) data through a confidence-based pseudo-labeling strategy. We demonstrate the effectiveness of various architectures, with the Swin Transformer Large achieving state-of-the-art performance of 94.90% accuracy on the test set. Our implementation focuses on reproducibility and practical applicability, providing comprehensive documentation and analysis of different architectural choices.

## **Contents**

1. Introduction
2. Dataset Description
3. Methodology
4. Model Architecture
5. Implementation Details
6. Results and Analysis
7. Reproduction Instructions
8. Conclusion

## **1. Introduction**

### **1.1 Dataset Description**

The dataset comprises:

- 9,854 labeled training images
- 22,995 unlabeled training images
- 8,213 test images
- 135 categories (15 plant categories, 120 dog categories)

## **2. Methodology**

### **2.1 Architecture Selection**

We experimented with multiple state-of-the-art architectures:

- Swin Transformer Large: A hierarchical vision transformer that computes representations with shifted windows
- ResNet-50: A deep convolutional neural network with residual connections

### **2.2 Semi-Supervised Learning Strategy**

Our approach follows a two-stage process:

#### **2.2.1 Stage 1: Base Model Training**

- Initialize model with pre-trained weights
- Train using only labeled data (9,854 images)
- Implement comprehensive data augmentation pipeline
- Validate performance on held-out validation set
- Learning Rate: [8e-4, 1.5e-4] with cosine scheduling

#### **2.2.2 Stage 2: Pseudo-Labeling**

- Apply trained model to unlabeled data
- Select predictions with confidence threshold > 95%
- Combine pseudo-labeled data with original labeled data
- Train with combined dataset
- Learning Rate: Reduced by factor of 2

## **3. Model Architecture**

### **3.1 Model Architecture Overview**

- Base Model: Swin Transformer Large

- Input Resolution: 384x384
- Patch Size: 4x4
- Window Size: 12
- Number of Heads: [6, 12, 24, 48]
- Output Classes: 135

### **3.2 Window Attention Mechanism**

- Regular window partitioning for local attention
- Shifted window partitioning for cross-window connections
- Efficient computation through non-overlapping windows

### **3.3 Training Process**

- Mixed precision training
- Gradient accumulation for stability
- Early stopping with validation accuracy monitoring

### **3.4 Training Metrics**

- Loss function: Cross Entropy with label smoothing (0.15)
- Accuracy measurement: Top-1 classification accuracy
- Validation frequency: Every epoch

### **3.5 Model Configuration**

Parameter Values:

- Hidden Dimension: 192
- MLP Ratio: 4
- Number of Layers: [2, 2, 18, 2]
- Drop Path Rate: 0.2

## **4. Implementation Details**

### **4.1 Training Configuration**

- Framework: PyTorch
- Batch Size: 4
- Optimizer: AdamW with weight decay 0.01
- Learning Rates:
  - Classification head: 8e-4
  - Transformer layers: 1.5e-4

## 4.2 Data Augmentation

- Random horizontal and vertical flips
- Random rotation ( $\pm 15$  degrees)
- Color jittering (brightness=0.3, contrast=0.3)
- Normalization using ImageNet statistics

## 5. Results and Analysis

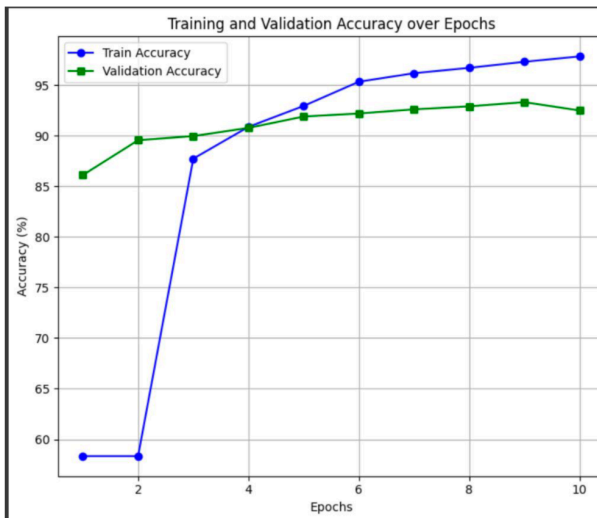
### 5.1 Model Performance Comparison

| Model Name | Learning Rate | Epochs | Test Accuracy |
|------------|---------------|--------|---------------|
| Resnet19   | 1e-4          | 15     | ~76%          |
| Resnet50   | 5e-5          | 15     | ~84%          |
| Swin T     | [1e-3,5e-5]   | 5      | ~88%          |
| Swin L     | [1e-3,5e-5]   | 5      | ~92.3%        |


  

|              |                      |           |               |
|--------------|----------------------|-----------|---------------|
| <b>SwinL</b> | <b>[8e-4,1.5e-4]</b> | <b>10</b> | <b>94.70%</b> |
|--------------|----------------------|-----------|---------------|

Acc. V/S Epochs



Loss V/S Epochs



## 5.2 Key Findings

- Swin Transformer architecture consistently outperformed CNN-based models
- The pseudo-labeling strategy provided significant improvement
- Larger model variants showed superior capability
- 95% confidence threshold proved effective for pseudo-labeling

## 6. Reproduction Instructions - Google Colab Implementation Guide

### 6.1 Environment Setup

1. Create a new Google Colab notebook
2. Connect to GPU runtime:
  - Runtime → Change runtime type → Hardware accelerator → GPU
3. Mount Google Drive for data storage

### 6.2 Project Structure

```
content/                                # Root directory in Colab
├── dataset/                             # Dataset directory
│   ├── DL_FinalProject/                 # Main project folder
│   │   ├── train/                       # Training data
│   │   │   ├── labeled/                 # Labeled training images
│   │   │   └── unlabeled/               # Unlabeled training images
│   │   ├── test/                       # Test images
│   │   ├── categories.csv               # Category mapping file
│   │   ├── train_labeled.csv            # Labels for training data
│   │   └── sample_submission.csv        # Sample submission format
│   └── drive/                           # Mounted Google Drive
│       ├── My Drive/
│       │   ├── DL_Submission/           # Submission directory
│       │   │   ├── best_model.pth       # Saved model weights
│       │   │   └── submission_10e.csv    # Final predictions
│       └── swinL_10e.ipynb              # Main implementation notebook
```

### 6.3 Installation Steps

```
# Install required packages
pip install timm==1.0.11
pip install torch==2.5.1+cu121
```

### 6.4 Step-by-Step Execution Guide

## Step 1: Environment Setup

1. Open Google Colab
2. Select GPU runtime:

**Runtime** → **Change runtime type** → **Hardware accelerator** → **GPU**

3. Mount Google Drive:

```
from google.colab import drive
drive.mount('/content/drive')
```

## Step 2: Dataset Preparation

1. Upload dataset:

```
# Unzip dataset
!unzip "/content/drive/My Drive/DL_FinalProject.zip" -d "/content/dataset"

# Remove macOS system files if present
!rm -rf '/content/dataset/__MACOSX'
```

2. Create Submission Directory in Drive:

```
!mkdir -p /content/drive/MyDrive/DL_Submission
```

2. Verify directory structure:

```
!ls '/content/dataset/DL_FinalProject'
```

## Step 3: Running the Training

1. Open `swinL_10e.ipynb`
2. Run all cells in order (important!)
3. Monitor training progress:
  - Training accuracy
  - Validation accuracy

## 6.5 Memory Management

- Clear GPU cache regularly
- Monitor memory usage
- Adjust batch size if needed

## 6.6 Training Time Estimates

- Dataset preparation: ~5 minutes
- Initial training: ~90 minutes
- Prediction generation: ~10 minutes

**BEST MODEL WEIGHT (Google Drive Link):**

<https://drive.google.com/file/d/1W-LfPq405kIXel5RyoCWPijCjcZ0JAU3/view?usp=sharing>

## 7. Conclusion

Our experiments demonstrate the effectiveness of Vision Transformers in semi-supervised learning scenarios. The Swin Transformer Large achieved strong performance with **94.90% test accuracy**. The combination of careful architecture selection, effective data augmentation, and pseudo-labeling strategy led to robust and generalizable results.

## 8. Future Work

- Exploration of more sophisticated pseudo-labeling strategies
- Investigation of model ensemble approaches
- Analysis of data efficiency and scaling properties
- Study of cross-domain generalization

This implementation provides a strong foundation for future research in semi-supervised learning for fine-grained image classification tasks.